Deliberate Attack

```
library(conflicted)

library(kableExtra)
```

```
## Warning in !is.null(rmarkdown::metadata$output) && rmarkdown::metadata$output
## %in% : 'length(x) = 2 > 1' in coercion to 'logical(1)'
```

```
library(knitr)
library(broom.helpers)
library(broom)
library(dtplyr)
library(furrr)
```

```
## Loading required package: future
```

```
library(arrow)
library(glue)
library(fs)
library(tidyverse)
```

```
## -- Attaching packages --------------------------------------- tidyverse 1.3.2 --
```

```
## v ggplot2 3.4.0      v purrr   1.0.1
## v tibble  3.1.8      v dplyr   1.0.10
## v tidyr   1.2.1      v stringr 1.5.0
## v readr   2.1.3      v forcats 0.5.2
```

```
conflict_prefer("filter", "dplyr")
```

```
## [conflicted] Will prefer dplyr::filter over any other package
```

```
source("./analysis/utils.R", local = knit_global())
set_theme()
```

```
write_bib(.packages(), glue("./analysis/packages.bib"))
```

```
## Warning in utils::citation(..., lib.loc = lib.loc): no date field in DESCRIPTION
## file of package 'kableExtra'
```

```
sessionInfo()
```

```
## R version 4.2.2 (2022-10-31)
## Platform: x86_64-apple-darwin17.0 (64-bit)
## Running under: macOS Big Sur ... 10.16
##
## Matrix products: default
## BLAS:   /Library/Frameworks/R.framework/Versions/4.2/Resources/lib/libRblas.0.dylib
## LAPACK: /Library/Frameworks/R.framework/Versions/4.2/Resources/lib/libRlapack.dylib
##
## locale:
## [1] en_US.UTF-8/en_US.UTF-8/en_US.UTF-8/C/en_US.UTF-8/en_US.UTF-8
```

```
##
## attached base packages:
## [1] stats     graphics  grDevices datasets  utils     methods   base
##
## other attached packages:
##  [1] forcats_0.5.2         stringr_1.5.0         dplyr_1.0.10
##  [4] purrr_1.0.1           readr_2.1.3           tidyr_1.2.1
##  [7] tibble_3.1.8          ggplot2_3.4.0         tidyverse_1.3.2
## [10] fs_1.5.2              glue_1.6.2            arrow_10.0.1
## [13] furrr_0.3.1           future_1.30.0         dtplyr_1.2.2
## [16] broom_1.0.2           broom.helpers_1.11.0  knitr_1.41
## [19] kableExtra_1.3.4.9000 conflicted_1.1.0
##
## loaded via a namespace (and not attached):
##  [1] httr_1.4.4          bit64_4.0.5         jsonlite_1.8.4
##  [4] viridisLite_0.4.1   here_1.0.1          modelr_0.1.10
##  [7] assertthat_0.2.1    renv_0.16.0         googlesheets4_1.0.1
## [10] cellranger_1.1.0    yaml_2.3.6          globals_0.16.2
## [13] pillar_1.8.1        backports_1.4.1     digest_0.6.31
## [16] rvest_1.0.3         colorspace_2.0-3    htmltools_0.5.4
## [19] pkgconfig_2.0.3     listenv_0.9.0       haven_2.5.1
## [22] scales_1.2.1        webshot_0.5.4       svglite_2.1.1
## [25] tzdb_0.3.0          timechange_0.2.0    googledrive_2.0.0
## [28] generics_0.1.3      ellipsis_0.3.2      withr_2.5.0
## [31] cachem_1.0.6        cli_3.6.0           crayon_1.5.2
## [34] readxl_1.4.1        magrittr_2.0.3      memoise_2.0.1
## [37] evaluate_0.19       fansi_1.0.3         parallelly_1.34.0
## [40] xml2_1.3.3          tools_4.2.2         data.table_1.14.6
## [43] hms_1.1.2           gargle_1.2.1        lifecycle_1.0.3
## [46] reprex_2.0.2        munsell_0.5.0       compiler_4.2.2
## [49] systemfonts_1.0.4   rlang_1.0.6         grid_4.2.2
## [52] rstudioapi_0.14     rmarkdown_2.19      gtable_0.3.1
## [55] codetools_0.2-18    DBI_1.1.3           R6_2.5.1
## [58] lubridate_1.9.0     fastmap_1.1.0       bit_4.0.5
## [61] utf8_1.2.2          rprojroot_2.0.3     stringi_1.7.12
## [64] parallel_4.2.2      vctrs_0.5.1         dbplyr_2.2.1
## [67] tidyselect_1.2.0    xfun_0.36
```

```
data_dir <- path(glue("./data/{params$simulation}/results"))

success_fnames <-
  dir_ls(data_dir, glob = glue("*{params$simulation}*_trend.csv"))

# every fname is a simulation
success_raw_data <- get_data(success_fnames, read_csv) |>
  glimpse()
```

```
## Rows: 480
## Columns: 14
## $ fname                  <chr> "./data/arbitrary/results/bbox_100_dist_100_re~
## $ num_iteration          <dbl> 200, 200, 200, 200, 200, 200, 200, 200, 200, 2~
## $ attack_count           <dbl> 100, 100, 100, 100, 100, 100, 100, 100, 100, 1~
## $ success_count          <dbl> 0, 0, 0, 16, 36, 1, 2, 4, 12, 7, 0, 0, 1, 27, ~
## $ vanish_count           <dbl> 0, 0, 0, 7, 27, 1, 2, 4, 11, 6, 0, 0, 1, 26, 2~
## $ mislabel_count         <dbl> 0, 0, 0, 9, 9, 0, 0, 0, 1, 1, 0, 0, 0, 1, 1, 0~
```

```
## $ mislabel_intended_count <dbl> 0, 0, 0, 9, 9, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
## $ sample_count            <dbl> 116, 116, 118, 118, 119, 116, 116, 118, 118, 1~
## $ model_name              <ord> Cascade R-CNN, Faster R-CNN, RetinaNet, SSD, Y~
## $ loss_target             <ord> Mislabeling, Mislabeling, Mislabeling, Mislabe~
## $ attack_bbox             <chr> "ground_truth", "ground_truth", "ground_truth"~
## $ perturb_fun             <chr> "perturb_inside", "perturb_inside", "perturb_i~
## $ arbitrary_bbox_length   <dbl> 100, 100, 100, 100, 100, 100, 100, 100, 100, 1~
## $ boundary_distance       <dbl> 100, 100, 100, 100, 100, 100, 100, 100, 100, 1~
```

```
# expand success per simulation into 1 and 0s per row
success_expanded_data <- success_raw_data |>
  rename(
    bbox_dist = boundary_distance,
    bbox_len = arbitrary_bbox_length
  ) |>
  rowwise() |>
  mutate(success = list(rep(0:1, times = c(attack_count - success_count, success_count)))) |>
  unnest_longer(success) |>
  glimpse()
```

```
## Rows: 48,000
## Columns: 15
## $ fname                   <chr> "./data/arbitrary/results/bbox_100_dist_100_re~
## $ num_iteration           <dbl> 200, 200, 200, 200, 200, 200, 200, 200, 200, 2~
## $ attack_count            <dbl> 100, 100, 100, 100, 100, 100, 100, 100, 100, 1~
## $ success_count           <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
## $ vanish_count            <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
## $ mislabel_count          <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
## $ mislabel_intended_count <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
## $ sample_count            <dbl> 116, 116, 116, 116, 116, 116, 116, 116, 116, 1~
## $ model_name              <ord> Cascade R-CNN, Cascade R-CNN, Cascade R-CNN, C~
## $ loss_target             <ord> Mislabeling, Mislabeling, Mislabeling, Mislabe~
## $ attack_bbox             <chr> "ground_truth", "ground_truth", "ground_truth"~
## $ perturb_fun             <chr> "perturb_inside", "perturb_inside", "perturb_i~
## $ bbox_len                <dbl> 100, 100, 100, 100, 100, 100, 100, 100, 100, 1~
## $ bbox_dist               <dbl> 100, 100, 100, 100, 100, 100, 100, 100, 100, 1~
## $ success                 <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
```

```
# check whether attack count equals experiment settings
stopifnot(all(success_raw_data$attack_count == 100))

reps <- success_raw_data |>
  count(model_name, loss_target, arbitrary_bbox_length, boundary_distance) |>
  glimpse()
```

```
## Rows: 240
## Columns: 5
## $ model_name            <ord> YOLOv3, YOLOv3, YOLOv3, YOLOv3, YOLOv3, YOLOv3, ~
## $ loss_target           <ord> Vanishing, Vanishing, Vanishing, Vanishing, Vani~
## $ arbitrary_bbox_length <dbl> 10, 10, 10, 10, 50, 50, 50, 50, 100, 100, 100, 1~
## $ boundary_distance     <dbl> 10, 50, 100, 200, 10, 50, 100, 200, 10, 50, 100,~
## $ n                     <int> 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, ~
```

```
stopifnot(unique(reps$n) == 2)
```

```
# control both
model <- partial(glm_model, predictor = "bbox_dist * bbox_len")
data <- success_expanded_data

reg_res <- get_tidied_reg(model, data, return_mod = TRUE)
```

```
## `summarise()` has grouped output by 'model_name', 'loss_target'. You can
## override using the `.groups` argument.
```

```
reg_est <- reg_res$tidied

ext_sig(reg_est, "neg", "bbox_dist")
```

```
## ----------bbox_dist----------
## Total 15 predictors:
## 15 (100%) significant;
## 15 (100%) neg
```

```
## # A tibble: 15 x 9
## # Groups:   model_name, loss_target [15]
##    model_name    loss_ta~1 term  estim~2 std.e~3 stati~4 p.value conf.~5 conf.~6
##    <ord>         <ord>     <chr>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
##  1 YOLOv3        Vanishing bbox~  -0.015   0.002   -7.68       0  -0.018  -0.011
##  2 YOLOv3        Mislabel~ bbox~  -0.015   0.002   -8.54       0  -0.018  -0.011
##  3 YOLOv3        Untarget~ bbox~  -0.021   0.003   -7.44       0  -0.026  -0.015
##  4 SSD           Vanishing bbox~  -0.015   0.002   -7.06       0  -0.019  -0.011
##  5 SSD           Mislabel~ bbox~  -0.018   0.002   -7.55       0  -0.023  -0.014
##  6 SSD           Untarget~ bbox~  -0.018   0.002   -7.74       0  -0.023  -0.014
##  7 RetinaNet     Vanishing bbox~  -0.045   0.006   -7.19       0  -0.058  -0.033
##  8 RetinaNet     Mislabel~ bbox~  -0.031   0.007   -4.24       0  -0.047  -0.018
##  9 RetinaNet     Untarget~ bbox~  -0.038   0.005   -7.45       0  -0.049  -0.029
## 10 Faster R-CNN  Vanishing bbox~  -0.061   0.01    -6.41       0  -0.081  -0.044
## 11 Faster R-CNN  Mislabel~ bbox~  -0.054   0.012   -4.66       0  -0.08   -0.034
## 12 Faster R-CNN  Untarget~ bbox~  -0.044   0.006   -8.01       0  -0.056  -0.034
## 13 Cascade R-CNN Vanishing bbox~  -0.063   0.01    -6.58       0  -0.083  -0.046
## 14 Cascade R-CNN Mislabel~ bbox~  -0.062   0.012   -5.24       0  -0.088  -0.041
## 15 Cascade R-CNN Untarget~ bbox~  -0.061   0.008   -7.54       0  -0.079  -0.047
## # ... with abbreviated variable names 1: loss_target, 2: estimate,
## #   3: std.error, 4: statistic, 5: conf.low, 6: conf.high
```

```
ext_sig(reg_est, "pos", "bbox_len")
```

```
## ----------bbox_len----------
## Total 15 predictors:
## 14 (93%) significant;
## 14 (93%) pos
```

```
## # A tibble: 14 x 9
## # Groups:   model_name, loss_target [14]
##    model_name    loss_ta~1 term  estim~2 std.e~3 stati~4 p.value conf.~5 conf.~6
##    <ord>         <ord>     <chr>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
##  1 YOLOv3        Vanishing bbox~   0.03    0.002   19.6        0   0.027   0.033
##  2 YOLOv3        Mislabel~ bbox~   0.019   0.001   16.6        0   0.016   0.021
##  3 YOLOv3        Untarget~ bbox~   0.007   0.001    6.53       0   0.005   0.009
##  4 SSD           Vanishing bbox~   0.024   0.001   17.7        0   0.021   0.027
##  5 SSD           Mislabel~ bbox~   0.02    0.001   16.0        0   0.017   0.022
```

```
##  6 SSD           Untarget~ bbox~   0.013   0.001   11.7        0   0.011   0.015
##  7 RetinaNet     Vanishing bbox~   0.016   0.001   10.6        0   0.013   0.019
##  8 RetinaNet     Mislabel~ bbox~   0.008   0.002    4.54       0   0.005   0.012
##  9 RetinaNet     Untarget~ bbox~   0.005   0.001    3.97       0   0.003   0.008
## 10 Faster R-CNN  Vanishing bbox~   0.011   0.001    7.13       0   0.008   0.014
## 11 Faster R-CNN  Mislabel~ bbox~   0.007   0.002    3.71       0   0.003   0.01
## 12 Faster R-CNN  Untarget~ bbox~   0.005   0.001    3.68       0   0.002   0.007
## 13 Cascade R-CNN Vanishing bbox~   0.015   0.002    9.40       0   0.012   0.018
## 14 Cascade R-CNN Mislabel~ bbox~   0.01    0.002    5.80       0   0.006   0.013
## # ... with abbreviated variable names 1: loss_target, 2: estimate,
## #   3: std.error, 4: statistic, 5: conf.low, 6: conf.high
```

```
ext_sig(reg_est, "both", "bbox_dist:bbox_len")
```

```
## ----------bbox_dist:bbox_len----------
## Total 15 predictors:
## 7 (47%) significant;
## 7 (47%) both
```

```
## # A tibble: 7 x 9
## # Groups:   model_name, loss_target [7]
##   model_name    loss_tar~1 term  estim~2 std.e~3 stati~4 p.value conf.~5 conf.~6
##   <ord>         <ord>      <chr>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
## 1 YOLOv3        Vanishing  bbox~       0       0   -6.08   0           0       0
## 2 SSD           Vanishing  bbox~       0       0   -4.82   0           0       0
## 3 SSD           Mislabeli~ bbox~       0       0   -2.46   0.014       0       0
## 4 RetinaNet     Vanishing  bbox~       0       0   -2.15   0.032       0       0
## 5 RetinaNet     Untargeted bbox~       0       0    6.92   0           0       0
## 6 Faster R-CNN  Untargeted bbox~       0       0    6.89   0           0       0
## 7 Cascade R-CNN Untargeted bbox~       0       0    6.20   0           0       0
## # ... with abbreviated variable names 1: loss_target, 2: estimate,
## #   3: std.error, 4: statistic, 5: conf.low, 6: conf.high
```

```
dist_lab <- "Perturb-Target Distance"
len_lab <- "Perturb Box Length"

pred_name <- glue("{dist_lab} and {len_lab}")
main_pt <- glue("longer {len_lab} or shorter {dist_lab} cause success rates to significantly increase f

print_statistics(reg_est, table_caption(pred_name, main_pt, "deliberate"))
```

Table 1: We run a logistic model regressing success against perturb-target distance and perturb box length in the deliberate attack experiment. Longer perturb box length or shorter perturb-target distance cause success rates to significantly increase for all model and attack combinations, except for perturb box length in untargeted attack on Cascade R-CNN. The interaction terms, even when significant, are negligibly close to 0. Table headers are explained in Appendix **??**.

Group	Regression							
Attack	term	sig	estimate	std.error	statistic	p.value	conf.low	conf.high
YOLOv3								
Vanishing	distance	*	-0.015	0.002	-7.681	0.000	-0.018	-0.011
	length	*	0.030	0.002	19.637	0.000	0.027	0.033
	distance * length	*	0.000	0.000	-6.081	0.000	0.000	0.000

	Mislabeling	distance	*	-0.015	0.002	-8.540	0.000	-0.018	-0.011
		length	*	0.019	0.001	16.603	0.000	0.016	0.021
		distance * length		0.000	0.000	-1.733	0.083	0.000	0.000
	Untargeted	distance	*	-0.021	0.003	-7.440	0.000	-0.026	-0.015
		length	*	0.007	0.001	6.528	0.000	0.005	0.009
		distance * length		0.000	0.000	1.467	0.142	0.000	0.000
SSD									
	Vanishing	distance	*	-0.015	0.002	-7.055	0.000	-0.019	-0.011
		length	*	0.024	0.001	17.747	0.000	0.021	0.027
		distance * length	*	0.000	0.000	-4.823	0.000	0.000	0.000
	Mislabeling	distance	*	-0.018	0.002	-7.553	0.000	-0.023	-0.014
		length	*	0.020	0.001	15.991	0.000	0.017	0.022
		distance * length	*	0.000	0.000	-2.458	0.014	0.000	0.000
	Untargeted	distance	*	-0.018	0.002	-7.742	0.000	-0.023	-0.014
		length	*	0.013	0.001	11.665	0.000	0.011	0.015
		distance * length		0.000	0.000	-0.873	0.383	0.000	0.000
RetinaNet									
	Vanishing	distance	*	-0.045	0.006	-7.187	0.000	-0.058	-0.033
		length	*	0.016	0.001	10.614	0.000	0.013	0.019
		distance * length	*	0.000	0.000	-2.147	0.032	0.000	0.000
	Mislabeling	distance	*	-0.031	0.007	-4.240	0.000	-0.047	-0.018
		length	*	0.008	0.002	4.541	0.000	0.005	0.012
		distance * length		0.000	0.000	-1.021	0.307	0.000	0.000
	Untargeted	distance	*	-0.038	0.005	-7.446	0.000	-0.049	-0.029
		length	*	0.005	0.001	3.969	0.000	0.003	0.008
		distance * length	*	0.000	0.000	6.925	0.000	0.000	0.000
Faster R-CNN									
	Vanishing	distance	*	-0.061	0.010	-6.407	0.000	-0.081	-0.044
		length	*	0.011	0.001	7.127	0.000	0.008	0.014
		distance * length		0.000	0.000	-0.490	0.624	0.000	0.000
	Mislabeling	distance	*	-0.054	0.012	-4.664	0.000	-0.080	-0.034
		length	*	0.007	0.002	3.706	0.000	0.003	0.010
		distance * length		0.000	0.000	-0.717	0.473	0.000	0.000
	Untargeted	distance	*	-0.044	0.006	-8.012	0.000	-0.056	-0.034
		length	*	0.005	0.001	3.676	0.000	0.002	0.007
		distance * length	*	0.000	0.000	6.889	0.000	0.000	0.000
Cascade R-CNN									
	Vanishing	distance	*	-0.063	0.010	-6.579	0.000	-0.083	-0.046
		length	*	0.015	0.002	9.395	0.000	0.012	0.018
		distance * length		0.000	0.000	-1.003	0.316	0.000	0.000
	Mislabeling	distance	*	-0.062	0.012	-5.240	0.000	-0.088	-0.041

	length	*	0.010	0.002	5.795	0.000	0.006	0.013
	distance * length		0.000	0.000	-0.122	0.903	0.000	0.000
Untargeted	distance	*	-0.061	0.008	-7.544	0.000	-0.079	-0.047
	length		0.002	0.001	1.498	0.134	-0.001	0.005
	distance * length	*	0.000	0.000	6.198	0.000	0.000	0.000

```
reg_mod <- reg_res$mod

newdata <- expand_grid(
  bbox_dist = linear_space(data$bbox_dist),
  bbox_len = unique(data$bbox_len)
) |>
  glimpse()
```

```
## Rows: 400
## Columns: 2
## $ bbox_dist <dbl> 10.00000, 10.00000, 10.00000, 10.00000, 11.91919, 11.91919, ~
## $ bbox_len  <dbl> 100, 10, 200, 50, 100, 10, 200, 50, 100, 10, 200, 50, 100, 1~
```

```
# type.predict = "link" by default
# https://broom.tidymodels.org/reference/augment.glm.html
# https://stackoverflow.com/questions/14423325/confidence-intervals-for-predictions-from-logistic-regre
reg_pred <- reg_mod |>
  summarize(augment(mod, newdata = newdata, se_fit = TRUE)) |>
  mutate(success = plogis(.fitted), ul = plogis(.fitted + 1.96 * .se.fit), ll = plogis(.fitted - 1.96 *
  glimpse()
```

```
## `summarise()` has grouped output by 'model_name', 'loss_target'. You can
## override using the `.groups` argument.
```

```
## Rows: 6,000
## Columns: 9
## Groups: model_name, loss_target [15]
## $ model_name  <ord> YOLOv3, YOLOv3, YOLOv3, YOLOv3, YOLOv3, YOLOv3, YOLOv3, YO~
## $ loss_target <ord> Vanishing, Vanishing, Vanishing, Vanishing, Vanishing, Van~
## $ bbox_dist   <dbl> 10.00000, 10.00000, 10.00000, 10.00000, 11.91919, 11.91919~
## $ bbox_len    <dbl> 100, 10, 200, 50, 100, 10, 200, 50, 100, 10, 200, 50, 100,~
## $ .fitted     <dbl> 1.1206648, -1.4569515, 3.9846829, -0.3113443, 1.0753688, -~
## $ .se.fit     <dbl> 0.08354473, 0.11366498, 0.19249889, 0.08011917, 0.08206361~
## $ success     <dbl> 0.7541120, 0.1889340, 0.9817412, 0.4227867, 0.7456166, 0.1~
## $ ul          <dbl> 0.7831999, 0.2254527, 0.9874075, 0.4614987, 0.7749042, 0.2~
## $ ll          <dbl> 0.7225041, 0.1571306, 0.9735935, 0.3850003, 0.7139250, 0.1~
```

```
arb_cap <- glue("{bold_tex('A deliberate attack obfuscates intent with increased success for all models

arb_cap
```

```
## \textbf{A deliberate attack obfuscates intent with increased success for all models and attacks:}  W
```

```
g <- success_expanded_data |> ggplot(aes(bbox_dist, success, color = bbox_len, group = bbox_len)) +
  stat_summary(fun.data = "mean_cl_boot") +
  facet_grid(cols = vars(model_name), rows = vars(loss_target))

# https://github.com/tidyverse/ggplot2/blob/ef00be7e2016e1259b4aef7f7c85651df123beff/R/geom-smooth.r#L1
g <- g + geom_ribbon(
```

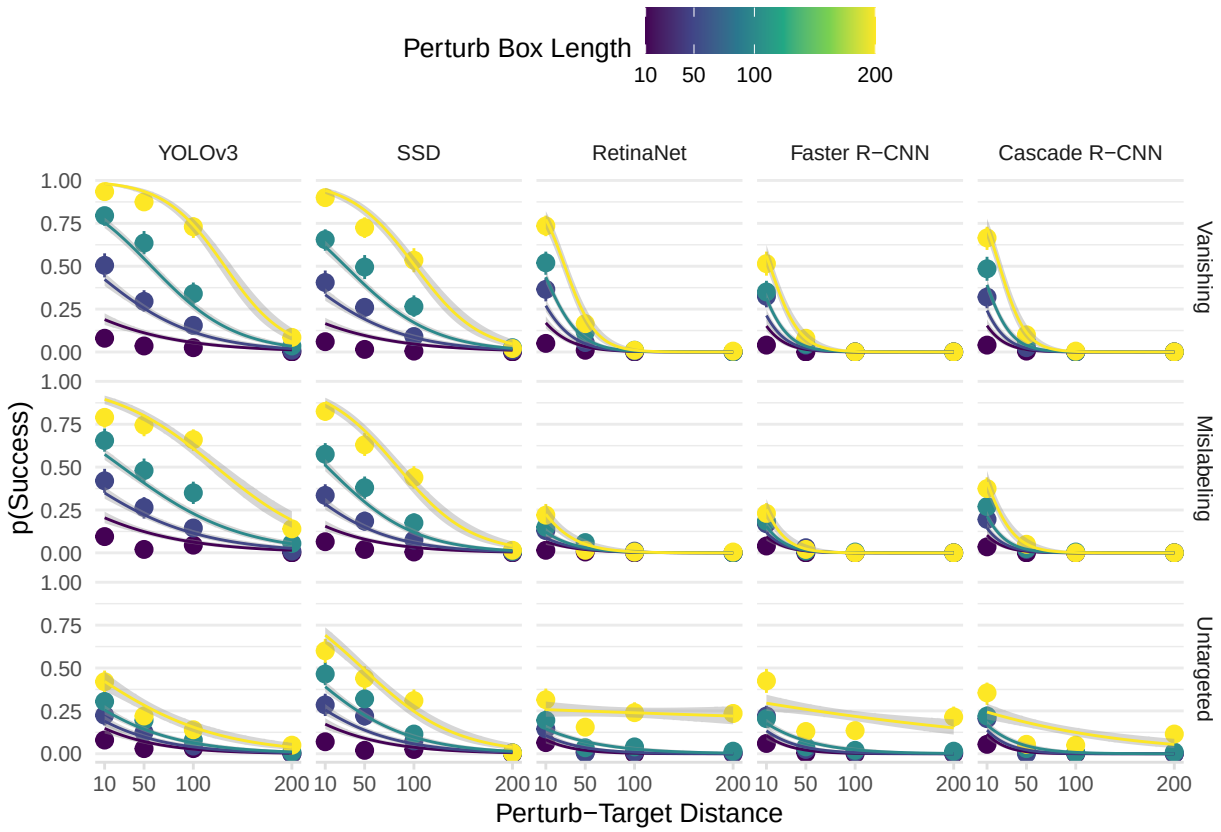

Figure 1: **A deliberate attack obfuscates intent with increased success for all models and attacks:** We implement intent obfuscating attack by perturbing an arbitrary non-overlapping square region to disrupt a randomly selected target object at various lengths and distances. The binned summaries and regression trendlines graph success proportion against perturb-target distance and perturb box length in the deliberate attack experiment. Errors are 95% confidence intervals. and every point aggregates success over 200 images. The deliberate attack multiplies success as compared to the randomized attack (Figure **??**), especially at close perturb-target distance and large perturb box length. Full details are given in Section **??**.

```
  data = reg_pred, aes(ymin = ll, ymax = ul),
  fill = "grey60", linetype = 0, alpha = 0.4
) +
  geom_line(data = reg_pred)

g + labs(x = dist_lab, y = "p(Success)") +
  scale_x_continuous(breaks = unique(success_expanded_data$bbox_dist)) +
  scale_color_viridis_c(name = len_lab, breaks = unique(success_expanded_data$bbox_len))
```

```
get_reg_vars <- function(data) {
  data |> select(bbox_dist, bbox_size_perturb, model_name, loss_target, success, object)
}

# run random.Rmd 1st
rand_dist_size <- readRDS("./analysis/rand_dist_size.RDS") |>
  mutate(object = 1) |>
  get_reg_vars() |>
  glimpse()
```

```
## Rows: 75,000
## Columns: 6
## $ bbox_dist         <dbl> 11.250000, 74.020000, 267.290000, 161.231650, 61.260~
## $ bbox_size_perturb <dbl> 4705.2345, 3803.1889, 595.2576, 29362.0050, 43664.54~
## $ model_name        <ord> Cascade R-CNN, Cascade R-CNN, Cascade R-CNN, Cascade~
## $ loss_target       <ord> Mislabeling, Mislabeling, Mislabeling, Mislabeling, ~
## $ success           <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
## $ object            <dbl> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1~
```

```
comb_dist_size <- success_expanded_data |>
  mutate(object = 0, bbox_size_perturb = bbox_len^2) |>
  get_reg_vars() |>
  bind_rows(rand_dist_size) |>
  mutate(
    bbox_dist = bbox_dist / 1e2,
    bbox_size_perturb = bbox_size_perturb / 1e5
  ) |>
  glimpse()
```

```
## Rows: 123,000
## Columns: 6
## $ bbox_dist         <dbl> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1~
## $ bbox_size_perturb <dbl> 0.1, 0.1, 0.1, 0.1, 0.1, 0.1, 0.1, 0.1, 0.1, 0.1, 0.~
## $ model_name        <ord> Cascade R-CNN, Cascade R-CNN, Cascade R-CNN, Cascade~
## $ loss_target       <ord> Mislabeling, Mislabeling, Mislabeling, Mislabeling, ~
## $ success           <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
## $ object            <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
```

```
stopifnot(nrow(comb_dist_size) == nrow(success_expanded_data) +
            nrow(rand_dist_size) && sum(is.na(comb_dist_size)) == 0)
```

```
# control both
model <- partial(glm_model, predictor = "object + bbox_dist * bbox_size_perturb")
data <- comb_dist_size

reg_est <- get_tidied_reg(model, data)
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## `summarise()` has grouped output by 'model_name', 'loss_target'. You can
## override using the `.groups` argument.
```

```
ext_sig(reg_est, "both", "object")
```

```
## ----------object----------
## Total 15 predictors:
## 10 (67%) significant;
## 10 (67%) both
```

```
## # A tibble: 10 x 9
## # Groups:   model_name, loss_target [10]
##    model_name    loss_ta~1 term  estim~2 std.e~3 stati~4 p.value conf.~5 conf.~6
##    <ord>         <ord>     <chr>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
##  1 YOLOv3        Vanishing obje~  -0.317   0.063   -5.03   0      -0.44   -0.193
##  2 YOLOv3        Untarget~ obje~  -0.201   0.079   -2.54   0.011  -0.357  -0.046
##  3 SSD           Untarget~ obje~   0.176   0.066    2.66   0.008   0.047   0.306
##  4 RetinaNet     Vanishing obje~  -0.551   0.093   -5.95   0      -0.734  -0.37
##  5 RetinaNet     Untarget~ obje~  -0.448   0.091   -4.90   0      -0.628  -0.269
##  6 Faster R-CNN  Vanishing obje~  -0.768   0.113   -6.81   0      -0.991  -0.549
##  7 Faster R-CNN  Mislabel~ obje~  -0.384   0.139   -2.77   0.006  -0.657  -0.113
##  8 Faster R-CNN  Untarget~ obje~  -0.275   0.089   -3.10   0.002  -0.449  -0.101
##  9 Cascade R-CNN Vanishing obje~  -0.665   0.104   -6.40   0      -0.87   -0.462
## 10 Cascade R-CNN Mislabel~ obje~  -0.282   0.117   -2.40   0.016  -0.513  -0.052
## # ... with abbreviated variable names 1: loss_target, 2: estimate,
## #   3: std.error, 4: statistic, 5: conf.low, 6: conf.high
```

```
ext_sig(reg_est, "neg", "bbox_dist")
```

```
## ----------bbox_dist----------
## Total 15 predictors:
## 15 (100%) significant;
## 15 (100%) neg
```

```
## # A tibble: 15 x 9
## # Groups:   model_name, loss_target [15]
##    model_name    loss_ta~1 term  estim~2 std.e~3 stati~4 p.value conf.~5 conf.~6
##    <ord>         <ord>     <chr>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
##  1 YOLOv3        Vanishing bbox~   -1.71   0.076   -22.4       0   -1.86   -1.56
##  2 YOLOv3        Mislabel~ bbox~   -1.52   0.066   -22.8       0   -1.65   -1.39
##  3 YOLOv3        Untarget~ bbox~   -1.97   0.114   -17.4       0   -2.20   -1.75
##  4 SSD           Vanishing bbox~   -1.92   0.084   -23.0       0   -2.09   -1.76
##  5 SSD           Mislabel~ bbox~   -1.95   0.091   -21.4       0   -2.13   -1.78
##  6 SSD           Untarget~ bbox~   -2.06   0.093   -22.0       0   -2.25   -1.88
```

```
##  7 RetinaNet     Vanishing bbox~   -4.95   0.261   -19.0       0   -5.47   -4.45
##  8 RetinaNet     Mislabel~ bbox~   -3.97   0.357   -11.1       0   -4.70   -3.30
##  9 RetinaNet     Untarget~ bbox~   -1.33   0.106   -12.6       0   -1.55   -1.13
## 10 Faster R-CNN  Vanishing bbox~   -6.00   0.408   -14.7       0   -6.83   -5.23
## 11 Faster R-CNN  Mislabel~ bbox~   -5.87   0.483   -12.1       0   -6.86   -4.96
## 12 Faster R-CNN  Untarget~ bbox~   -1.80   0.124   -14.6       0   -2.05   -1.57
## 13 Cascade R-CNN Vanishing bbox~   -6.50   0.388   -16.7       0   -7.28   -5.76
## 14 Cascade R-CNN Mislabel~ bbox~   -6.32   0.438   -14.4       0   -7.21   -5.49
## 15 Cascade R-CNN Untarget~ bbox~   -2.46   0.159   -15.5       0   -2.79   -2.16
## # ... with abbreviated variable names 1: loss_target, 2: estimate,
## #   3: std.error, 4: statistic, 5: conf.low, 6: conf.high
```

```
ext_sig(reg_est, "pos", "bbox_size_perturb")
```

```
## ----------bbox_size_perturb----------
## Total 15 predictors:
## 14 (93%) significant;
## 14 (93%) pos
```

```
## # A tibble: 14 x 9
## # Groups:   model_name, loss_target [14]
##    model_name    loss_ta~1 term  estim~2 std.e~3 stati~4 p.value conf.~5 conf.~6
##    <ord>         <ord>     <chr>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
##  1 YOLOv3        Vanishing bbox~   8.59    0.367   23.4        0   7.88     9.32
##  2 YOLOv3        Mislabel~ bbox~   4.54    0.253   17.9        0   4.05     5.04
##  3 YOLOv3        Untarget~ bbox~   1.59    0.17     9.36       0   1.26     1.93
##  4 SSD           Vanishing bbox~   5.88    0.296   19.9        0   5.31     6.47
##  5 SSD           Mislabel~ bbox~   4.23    0.249   17.0        0   3.75     4.73
##  6 SSD           Untarget~ bbox~   1.96    0.187   10.5        0   1.60     2.33
##  7 RetinaNet     Vanishing bbox~   2.69    0.251   10.7        0   2.21     3.19
##  8 RetinaNet     Mislabel~ bbox~   1.16    0.231    5.03       0   0.712    1.62
##  9 RetinaNet     Untarget~ bbox~   1.68    0.165   10.2        0   1.36     2.00
## 10 Faster R-CNN  Vanishing bbox~   2.06    0.256    8.05       0   1.57     2.58
## 11 Faster R-CNN  Untarget~ bbox~   2.10    0.182   11.6        0   1.75     2.46
## 12 Cascade R-CNN Vanishing bbox~   2.90    0.277   10.5        0   2.38     3.46
## 13 Cascade R-CNN Mislabel~ bbox~   0.886   0.22     4.02       0   0.459    1.32
## 14 Cascade R-CNN Untarget~ bbox~   0.913   0.161    5.68       0   0.598    1.23
## # ... with abbreviated variable names 1: loss_target, 2: estimate,
## #   3: std.error, 4: statistic, 5: conf.low, 6: conf.high
```

```
ext_sig(reg_est, "both", "bbox_dist:bbox_size_perturb")
```

```
## ----------bbox_dist:bbox_size_perturb----------
## Total 15 predictors:
## 8 (53%) significant;
## 8 (53%) both
```

```
## # A tibble: 8 x 9
## # Groups:   model_name, loss_target [8]
##   model_name    loss_tar~1 term  estim~2 std.e~3 stati~4 p.value conf.~5 conf.~6
##   <ord>         <ord>      <chr>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
## 1 YOLOv3        Vanishing  bbox~  -3.26    0.31   -10.5    0       -3.87  -2.66
## 2 YOLOv3        Mislabeli~ bbox~  -0.908   0.231   -3.94   0       -1.36  -0.462
## 3 SSD           Vanishing  bbox~  -2.26    0.289   -7.83   0       -2.84  -1.71
## 4 SSD           Mislabeli~ bbox~  -1.51    0.282   -5.36   0       -2.08  -0.971
## 5 RetinaNet     Untargeted bbox~   1.77    0.203    8.70   0        1.37   2.17
```

```
## 6 Faster R-CNN  Mislabeli~ bbox~   2.06    0.747    2.75   0.006    0.44   3.36
## 7 Faster R-CNN  Untargeted bbox~   2.23    0.233    9.57   0         1.78   2.69
## 8 Cascade R-CNN Untargeted bbox~   2.09    0.216    9.69   0         1.67   2.52
## # ... with abbreviated variable names 1: loss_target, 2: estimate,
## #   3: std.error, 4: statistic, 5: conf.low, 6: conf.high
```

```
dist_lab <- "Perturb-Target Distance (100 pixels)"
size_lab <- "Perturb Box Size (100,000 squared pixels)"

pred_name <- glue("object (versus non-object), with {dist_lab} and {size_lab} as covariates")
main_pt <- "perturbing an object (in the randomized attack) rather than a non-object (in the deliberate

tab_cap <- glue("We combined the data in the randomized and deliberate attack experiments to run a logi

print_statistics(reg_est, tab_cap)
```

Table 2: We combined the data in the randomized and deliberate attack experiments to run a logistic model regressing success against object (versus non-object), with perturb-target distance (100 pixels) and perturb box size (100,000 squared pixels) as covariates. The "object" term codes object as 1 and non-object as 0. Perturbing an object (in the randomized attack) rather than a non-object (in the deliberate attack) significantly decreases success rates for most model and attack combinations, after controlling for perturb sizes and perturb-target distances. Table headers are explained in Appendix **??**.

Group		Regression						
Attack	term	sig	estimate	std.error	statistic	p.value	conf.low	conf.high
YOLOv3								
Vanishing	object	*	-0.317	0.063	-5.031	0.000	-0.440	-0.193
	distance	*	-1.711	0.076	-22.383	0.000	-1.863	-1.563
	size	*	8.585	0.367	23.423	0.000	7.878	9.315
	distance * size	*	-3.258	0.310	-10.498	0.000	-3.872	-2.655
Mislabeling	object		-0.026	0.059	-0.440	0.660	-0.141	0.089
	distance	*	-1.515	0.066	-22.796	0.000	-1.647	-1.386
	size	*	4.538	0.253	17.940	0.000	4.050	5.041
	distance * size	*	-0.908	0.231	-3.938	0.000	-1.365	-0.462
Untargeted	object	*	-0.201	0.079	-2.544	0.011	-0.357	-0.046
	distance	*	-1.970	0.114	-17.351	0.000	-2.197	-1.752
	size	*	1.593	0.170	9.364	0.000	1.265	1.932
	distance * size		0.356	0.250	1.423	0.155	-0.149	0.837
SSD								
Vanishing	object		0.096	0.063	1.532	0.125	-0.027	0.219
	distance	*	-1.924	0.084	-22.955	0.000	-2.090	-1.762
	size	*	5.883	0.296	19.896	0.000	5.313	6.472
	distance * size	*	-2.263	0.289	-7.826	0.000	-2.844	-1.707
Mislabeling	object		-0.039	0.064	-0.609	0.542	-0.166	0.087
	distance	*	-1.953	0.091	-21.407	0.000	-2.134	-1.776
	size	*	4.228	0.249	16.958	0.000	3.749	4.726

	distance * size	*	-1.509	0.282	-5.356	0.000	-2.077	-0.971
Untargeted	object	*	0.176	0.066	2.661	0.008	0.047	0.306
	distance	*	-2.060	0.093	-22.041	0.000	-2.246	-1.879
	size	*	1.958	0.187	10.482	0.000	1.599	2.331
	distance * size		-0.227	0.244	-0.929	0.353	-0.719	0.240
RetinaNet								
Vanishing	object	*	-0.551	0.093	-5.947	0.000	-0.734	-0.370
	distance	*	-4.949	0.261	-18.960	0.000	-5.472	-4.448
	size	*	2.686	0.251	10.722	0.000	2.208	3.190
	distance * size		-0.881	0.569	-1.548	0.122	-2.035	0.197
Mislabeling	object		-0.245	0.136	-1.799	0.072	-0.513	0.022
	distance	*	-3.968	0.357	-11.109	0.000	-4.697	-3.297
	size	*	1.163	0.231	5.032	0.000	0.712	1.621
	distance * size		0.117	0.696	0.168	0.867	-1.323	1.403
Untargeted	object	*	-0.448	0.091	-4.902	0.000	-0.628	-0.269
	distance	*	-1.333	0.106	-12.560	0.000	-1.546	-1.130
	size	*	1.675	0.165	10.157	0.000	1.355	2.002
	distance * size	*	1.766	0.203	8.701	0.000	1.373	2.170
Faster R-CNN								
Vanishing	object	*	-0.768	0.113	-6.813	0.000	-0.991	-0.549
	distance	*	-6.002	0.408	-14.728	0.000	-6.829	-5.230
	size	*	2.062	0.256	8.052	0.000	1.572	2.577
	distance * size		-1.190	0.905	-1.315	0.188	-3.059	0.485
Mislabeling	object	*	-0.384	0.139	-2.770	0.006	-0.657	-0.113
	distance	*	-5.868	0.483	-12.144	0.000	-6.858	-4.961
	size		0.461	0.252	1.832	0.067	-0.029	0.958
	distance * size	*	2.055	0.747	2.752	0.006	0.440	3.362
Untargeted	object	*	-0.275	0.089	-3.096	0.002	-0.449	-0.101
	distance	*	-1.804	0.124	-14.599	0.000	-2.053	-1.568
	size	*	2.104	0.182	11.585	0.000	1.752	2.464
	distance * size	*	2.226	0.233	9.570	0.000	1.778	2.690
Cascade R-CNN								
Vanishing	object	*	-0.665	0.104	-6.395	0.000	-0.870	-0.462
	distance	*	-6.496	0.388	-16.731	0.000	-7.279	-5.757
	size	*	2.905	0.277	10.474	0.000	2.378	3.465
	distance * size		-1.579	0.840	-1.881	0.060	-3.310	-0.020
Mislabeling	object	*	-0.282	0.117	-2.402	0.016	-0.513	-0.052
	distance	*	-6.317	0.438	-14.410	0.000	-7.210	-5.489
	size	*	0.886	0.220	4.018	0.000	0.459	1.325
	distance * size		1.310	0.746	1.757	0.079	-0.265	2.666
Untargeted	object		-0.175	0.100	-1.739	0.082	-0.371	0.022

distance	*	-2.464	0.159	-15.457	0.000	-2.786	-2.160
size	*	0.913	0.161	5.677	0.000	0.598	1.229
distance * size	*	2.093	0.216	9.686	0.000	1.670	2.519